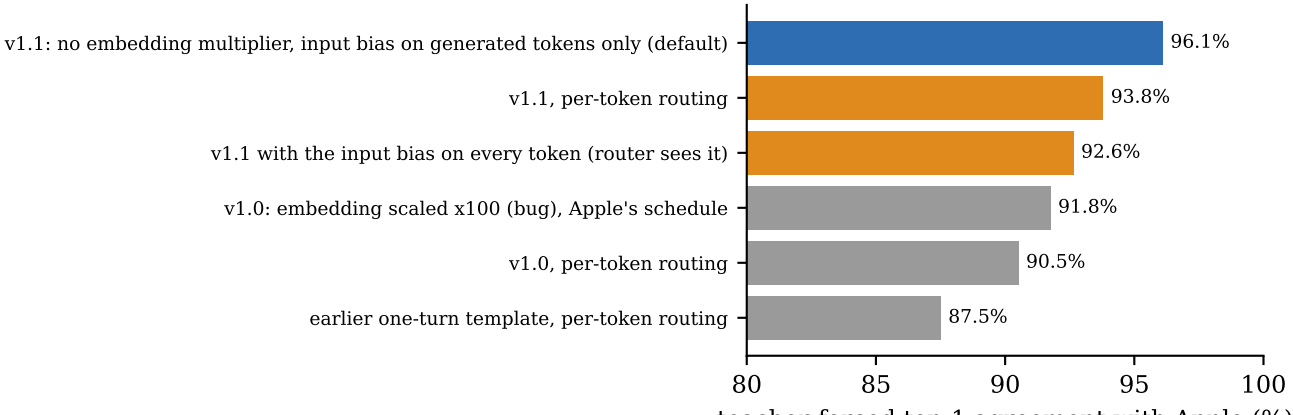
v1.1: no embedding multiplier, input bias on generated tokens only (default)
96.1%
v1.1, per-token routing
93.8%
v1.1 with the input bias on every token (router sees it)
92.6%
v1.0: embedding scaled x100 (bug), Apple's schedule
91.8%
v1.0, per-token routing
90.5%
earlier one-turn template, per-token routing
87.5%
80
85
90
95
100
teacher-forced top-1 agreement with Apple (%)